

# Spectral learning for structured partially observable environments

Lucas Langer

McGill University

*lucas.langer@mail.mcgill.ca*

August 18, 2015

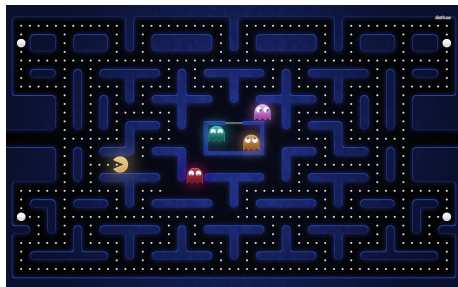
# Overview

- 1 Predictive state representations
- 2 The Base System
- 3 Experimental results
- 4 Learning the Base System

# Structure partially observable environments

## Structured Environments

- ① Goal: Predictions
- ② Plan: Exploit structure
- ③ Example: Pacman



# Structure partially observable environments

- Goal: Improve predictive performance
- Example: Pacman
- Plan: Represent structure in model

# PSRs: The Timing Case

- PSRs defined by:  $\langle \alpha_0, \{A_\sigma\}, \alpha_\infty \rangle$   
 $\alpha_0$ : Initial weighting on states  $1 \times n$   
 $A_\sigma$  Transition matrix  $n \times n$   
 $\alpha_\infty$ : Normalizer  $n \times 1$
- PSRs compute probabilities of observations  
 $\sigma$  represents one time unit,  $\sigma^k$  represents  $k$  time units  
 $f(\sigma^k) = \alpha_0 * A_\sigma^k * \alpha_\infty$
- Examples of a PSRs: HMMs, POMDPs

# Overview of Learning


- Spectral algorithms can learn PSRs
- Flavour of learning algorithm:
  - Step 1: Represent data as a matrix
  - Step 2: Singular value decomposition
  - Step 3: Pick number of states for PSR
  - Step 4: Learn the PSR with matrix computations

# The Base System

- Number representations:  $11 = 2^3 + 2^1 + 2^0$
- Timing queries  $f(a^11) = \alpha * A(a^8) * A(a^2) * A(a^1)$
- Motivation:
  - 1) Express transitions directly to avoid error build up
  - 2) Faster queries. Discussion  $\alpha_0 * (A_\sigma)^k$

# Timing with the Base



Agent goes through loops until leaving through an exit state. Exit states have transition probabilities of 0.4 and 0.6. Loop lengths are 64 and 16.






# Double Loop Results

$$||f_A - f_{A\text{Bar}}|| = (\sum (f_A(x) - f_{A\text{Bar}})^2)^{0.5}$$




# PacMan Labyrinth Results

$$||f_A - f_{ABar}|| = (\sum (f_A(x) - f_{ABar})^2)^{0.5}$$




# Distance Predictions

We use  $\alpha_0 * A_{\sigma}^k$  as a representation of state. Linear regression gives us a distance weighting on states.



# Wall Color Predictions

We paint the first loop green and the second loop red.



# Picking the Base System

- How do we pick transition operators?

Observations:  $\{a^{30}:10, a^{60}:5, b^{18}:15\}$

Desired Base System:  $A_a^{30}, A_b^{18}, A_a, A_b$

- Substring properties: **long, frequent, diverse**
- Solution: iterative greedy heuristic

# Computing with the Base System

- Using the Base System well involves requires good **string partitions**

Query string: "abcacb", Base System =  $\{A_{ab}, A_{bca}, A_{cb}, A_a, A_b\}$

Desired partition: "a—bca—cb""

- Goal: minimize matrices used
- Solution: dynamic programming

# Conclusion and Future Work

- What's left for the Base System?
  - 1) Theoretical analysis
  - 2) Test heuristics for selection and querying
  - 3) Further optimize heuristics

# Questions? Comments?